

Agent Harnesses for Evidence-Grounded Financial Prediction: Qualification-Gated Authority and Conditional Routing

Garros Gong

Department of Management Sciences

University of Waterloo

Waterloo, ON N2L 3G1 Canada

g7gong@uwaterloo.ca

Abstract

Hybrid financial forecasts need a rule for assigning forecast authority to a task specialist and updating the record after a decision changes. This paper develops a decision harness that conditions forecast authority on frozen task-specific qualification and current-case applicability, then reconstructs the complete record after selective category routing. The qualification audits grant forecast authority to numerical specialists in futures positioning and gasoline inflation. Other candidates retain the primary path when qualification fails or evidence is insufficient. Post-earnings routing produces imprecise accuracy changes that vary across the evaluated language-model backbones. Complete-record finalization preserves point-class consistency in the observed category changes. The evaluations show task-specific qualification outcomes and structural consistency after routing, with uncertain accuracy effects across the tested backbones.

1 Introduction

Evidence-grounded financial forecasts use an information cutoff to limit inputs to material available at decision time (Dawid, 1984; Jin et al., 2021). Each complete forecast record contains a point estimate, a class or ranking when relevant, any predictive interval, and cited evidence. When a large language model (LLM) and a task-specific statistical specialist disagree, forecast authority determines which candidate sets that record. Forecast combination pools competing predictions, but leaves the authority decision for a particular case open (Bates & Granger, 1969; Clemen, 1989). Repeated generations reveal factual inconsistency and uncertainty across meanings, making

disagreement observable before a decision is issued (Manakul et al., 2023; Farquhar et al., 2024). Selective classification formalizes abstention when an automated decision lacks sufficient support (Chow, 1970; Geifman & El-Yaniv, 2017). When another predictor has relevant expertise, learning to defer extends abstention to the decision of assigning the prediction to that expert (Madras et al., 2018; Mozannar & Sontag, 2020). Specialist competence still requires historical evidence for the target and information set at hand. The resulting question is what evidence qualifies a specialist to receive authority at a fixed cutoff and how an authority transfer should update the complete forecast record.

We develop a decision harness that freezes eligible evidence and governs which forecaster controls a complete forecast record.¹ The evaluation spans financial prediction tasks, including futures positioning, inflation nowcasting, and post-earnings stock reactions. The main contribution of this paper is a qualification-gated framework for assigning forecast authority. Historical task-specific performance determines specialist qualification, and current-case applicability limits numerical authority to the validated target and inputs. The qualification audit grants that authority to two numerical specialists and retains the primary path when other candidates fail or lack sufficient evidence. The secondary contribution couples conditional category routing to complete-record finalization, carrying a selected class through the linked point estimate, interval, and cited evidence. Post-earnings routing provides a model-sensitive test of that intervention, with the routing policy evaluated separately from the qualified numerical specialists. The rest of the paper examines how qualification evidence bears on case-level authority, then follows the consequences of an authority transfer through the complete forecast record.

2 Data and Decision Design

The forecast unit is a financial entity or event observed at an information cutoff for a specified target and horizon. Documents published or accepted by that cutoff form its evidence set, and realized outcomes enter later evaluation (Dawid, 1984; Jin et al., 2021). Specialist qualification requires frozen historical performance to meet a task-specific gate against a named control. Applicability checks whether the target, inputs, and source timing match that validated setting. Table 1 summarizes the three main outcomes, their sample counts, and their source boundaries. Secondary audits assess qualification and evidence sufficiency for additional candidate specialists (Table 2).

¹Code, configuration files, and frozen replay examples are available at <https://github.com/garroshub/AgentDecisionHarness>.

2.1 Data and Information Cutoffs

Weekly Commitments of Traders (COT) reports from the Commodity Futures Trading Commission (CFTC) anchor 104 forecast origins in 2013–2014 (Commodity Futures Trading Commission, 2026). Each year contributes 52 origins, and only reports available at each forecast origin enter the evidence set. The five-report horizon was fixed before evaluation, producing overlapping outcome windows across adjacent weekly origins, and the frozen positioning control extrapolates recent net-position momentum. The gasoline block contains 20 months excluded from specialist development, and the specialist’s training observations end in 2016. Frozen Archival Federal Reserve Economic Data (AL-FRED) Consumer Price Index (CPI) vintages and U.S. Energy Information Administration (EIA) gasoline history provide the pre-cutoff inputs (Croushore & Stark, 2001; U.S. Energy Information Administration, 2026). The outcome is the first-published Consumer Price Index for All Urban Consumers (CPI-U) gasoline value from the U.S. Bureau of Labor Statistics (BLS) (U.S. Bureau of Labor Statistics, 2025). The previously fixed baseline rule supplies the frozen reference proxy.

The post-earnings block uses a fixed issuer roster and an August–December 2025 window to obtain 30 events, one per issuer. Item 2.02 of a U.S. Securities and Exchange Commission (SEC) Form 8-K covers results of operations and financial condition (U.S. Securities and Exchange Commission, 2025). The earliest eligible filing per issuer must reach the SEC at or after 4:00 p.m. Eastern, the regular-session close (New York Stock Exchange, 2026). The event filing sets the cutoff timing, and each forecast packet draws from the latest earlier Form 10-K or 10-Q. The earlier Form 10-K or 10-Q is anonymized in full before a fixed keyword-weighted retriever selects up to 14 windows of 1,200 characters with a 900-character step for the LLM packet. Retrieval limits each LLM call to a compact, inspectable source context (Lewis et al., 2020). The frozen auxiliary text classifier uses that full, strictly anonymized filing. The Gemini 3.8 Flash model card reports a March 2026 knowledge cutoff, with some domains retaining information only through January 2025 (Google DeepMind, 2026). Search, grounding, and external tools remain disabled throughout the run. The evaluation locks all packet hashes and 30 prediction records before retrieving adjusted closing prices for the outcomes (Yahoo Finance, 2026). The primary run keeps realized outcomes sealed until all prediction records are locked. Later model checks reuse the same events after outcome release as sensitivity analyses.

2.2 Financial Outcomes and Forecast Variables

Noncommercial net positioning equals long minus short futures contracts, and open interest counts outstanding contracts at the origin report. The five-report positioning target divides the subsequent change in net contracts by origin-date open interest, producing a percent comparable across markets. Predicted changes are then ranked from the largest increase toward net-long positioning. Gasoline uses the first-published, seasonally adjusted month-over-month percentage change in the CPI-U gasoline component, fixing the evaluation outcome before subsequent revisions (U.S. Bureau of Labor Statistics, 2025).

Post-earnings returns use the issuer’s event-day and next-trading-day adjusted closes. The outcome subtracts the matching return on the SPDR S&P 500 ETF Trust (SPY), an exchange-traded fund tracking the S&P 500 (Brown & Warner, 1985; MacKinlay, 1997). The prespecified classes are positive above +1 percentage point, negative below -1 percentage point, and flat within that band. The one-point band separates small market-adjusted moves from directional reactions and was fixed before outcomes were assembled. Each forecast record contains a point estimate, a class or ranking when relevant, a nominal 90% interval, and cited pre-cutoff evidence. The output contract fixes the nominal interval level, and empirical coverage is evaluated separately. Ranking quality equals $(1 + \rho_S)/2$, where ρ_S is the Spearman correlation between predicted and realized market ranks at each origin. Gasoline point error is mean absolute error (MAE), which preserves the target’s percentage-point units (Hyndman & Koehler, 2006). Post-earnings classification accuracy is the share of issuer-events assigned the correct class. Empirical coverage checks the nominal interval level, and a proper interval score penalizes excess width and misses (Gneiting et al., 2007; Kuleshov et al., 2018). The 90% score is $IS_{0.10} = (U - L) + 20(L - Y)_+ + 20(Y - U)_+$, where L and U are the interval bounds and Y is the realized outcome (Gneiting & Raftery, 2007).

2.3 Qualification-Gated Decision Design

Specialist authority starts with qualification against frozen evidence from the same financial target. The active COT specialist is a standardized Ridge model with $\alpha = 10$ using crowding and one-, two-, four-, and eight-report momentum features. Its point gate requires mean rank quality above 0.5 overall and in both calendar years. Qualification also requires at least four of the five every-fifth-week offsets to clear 0.5, with combined rank quality matching or exceeding the four-report momentum control. The frozen interval composite is $0.70Q - 0.30|\text{Cov} - 0.90|$, where Q is rank

Table 1: Forecast Variables, Sample Descriptions, and Data Sources

Outcome variable	Forecast unit and sample	Sample summary	Pre-cutoff inputs and outcome source
Five-report change in noncommercial net futures positions, divided by origin-date open interest (%)	Futures market at a weekly origin; 104 ranking origins in 2013–2014	52 origins in each year	Available CFTC COT reports; the report five releases later supplies the outcome
First-published monthly CPI-U gasoline change, seasonally adjusted (%)	Gasoline component-month; 20 months unused in specialist development	20 evaluation months; training observations end in 2016	Frozen ALFRED CPI vintages and EIA gasoline history; initial BLS release supplies the outcome
Next-trading-day issuer return minus SPY return (percentage points)	Issuer-event; 30 events in August–December 2025	Mean -1.81 , standard deviation (SD) 11.38 ; negative 18, flat 5, positive 7	Pre-event SEC filings; frozen Yahoo Finance adjusted closing prices supply both returns

Notes: The COT count refers to ranking origins, each containing a market roster. The 20 gasoline months included in evaluation need not be consecutive. The CPI comparator is a frozen reference proxy from a previously fixed baseline rule.


quality and Cov is empirical coverage. Its weights prioritize market ranking and penalize departure from the nominal 90% coverage target. The interval stage requires nonnegative candidate-minus-control differences in that composite overall and by year. It also requires a proper interval score at or below the control. Gasoline uses a pooled Ridge model with $\alpha = 1$ trained through 2016. Qualification compares its MAE with the frozen reference proxy on the 20 months excluded from development. Passing that gate grants forecast authority for the gasoline component. For each numerical task, its frozen control remains in place outside the validated target, inputs, or source timing. The reproducibility index links frozen designs and executed checks to the canonical result manifest.²

Post-earnings routing starts after the event, anonymization, and packet audits have established case eligibility. Three independent LLM calls produce complete candidate records for each eligible event. The three-draw majority baseline (H3) is the representative complete record from a unique majority among structurally valid candidates. The first valid record supplies the fallback when no unique majority exists. Auxiliary routing is eligible only when all three candidate records are valid, their labels form an exact two-to-one split, and the frozen auxiliary label differs from the majority. The auxiliary classifier is a word term frequency-inverse document frequency (TF-IDF) logistic model trained through 2015 and applied to the full strictly anonymized filing. Each LLM call receives the fixed retrieval packet selected from that same filing. Figure 1 separates qualified

²The paper-facing evidence map appears in `paper/reproducibility/README.md`, with result paths and hashes in `paper/results/CANONICAL_RESULTS_MANIFEST.json`.

numerical ownership from the post-earnings disagreement trigger and follows both paths to record checking.³

Figure 1: Qualification-Gated Decision Harness



Without the trigger, H3 is retained. Quote matching is audited separately.

Notes: The upper path grants numerical authority after historical qualification and current-case applicability checks. The lower path shows the evaluated post-earnings trigger: three valid candidate records, an exact two-to-one class split, and an auxiliary label different from the majority. When the trigger is absent, the complete H3 record remains in place. Each changed class is finalized as a complete record before structural validation. Exact quote matching against the frozen packet is a separate provenance diagnostic.

When auxiliary authority is triggered, constrained finalization fixes the selected class and regenerates the point estimate, interval, and cited evidence as one record. Structural validity requires an allowed class, finite numerical fields, an ordered interval containing the point estimate, point–class agreement at the one-percentage-point boundaries, and a nonempty evidence quote. Exact quote matching against the frozen packet is recorded separately as a provenance diagnostic. If finalization fails or returns a different class, an existing valid candidate with the auxiliary label supplies the routed record when available. If that check also fails, the H3 record remains the complete-record fallback. The class-only override changes the H3 label without rebuilding the remaining fields and is retained solely for the record-integrity comparison.

³The written protocol initially described the disagreement case as minority support. The executable rule fixed before outcome release permitted a valid third class as well. The reproducibility record documents this discrepancy between the written and executable routing rules.

3 Results

Numerical specialists pass task-specific qualification gates, and post-earnings routing yields different observed accuracy changes across the evaluated backbones. Complete-record finalization carries a changed class into a structurally consistent forecast record.

3.1 Qualification-Gated Forecast Authority

The COT and gasoline entries in Table 2 provide the two positive qualification cases. The COT specialist clears both point and interval gates in the frozen confirmation block. The specialist achieves mean rank quality of 0.568 against 0.433 for the momentum control. Descriptive reanalysis gives year-specific values of 0.616 in 2013 and 0.520 in 2014, with all five nonoverlapping offsets above 0.5. The specialist also achieves a proper interval score of 42.53 against 61.49 for the control. The gasoline specialist clears its frozen holdout gate on the unused monthly block. Over the 20 months excluded from development, MAE is 2.252 for the Ridge model and 2.701 for the frozen reference proxy, a 16.6% reduction.

The secondary audits in Table 2 leave task-wide forecast authority unchanged. Frozen confirmation fails for the earnings-per-share (EPS) growth upgrade, whose mean quality-score change is -0.016 , and the macro-revision candidate reaches 50.0% classification accuracy. The Treasury exercise contributes interval evidence, with 89.3% leave-one-tenor-out coverage. The existing auction point forecast remains in place because this check does not qualify a replacement point forecast. Validation evidence is insufficient for the yield-curve and issuer-credit candidates, which rely on a surrogate stress test and two historical retrieval cases, respectively. The retrospective EPS year-over-year replay is neutral: all 27 three-draw predictions are unanimous, and H3 and the selective harness each reach 70.4% classification accuracy.

3.2 Backbone Dependence of Routing Authority

Table 3 compares the same post-earnings routing rule under three language-model backbones. The 30 outcomes include 18 negative, five flat, and seven positive reactions, so a constant-negative forecast reaches 60.0% accuracy. In the primary outcome-sealed Gemini 3.8 Flash run, routing raises accuracy from 26.7% to 33.3%, a gain of 6.7 percentage points. The case-bootstrap 95% interval spans $[-6.7, 20.0]$ percentage points, leaving the direction of the accuracy change uncertain. Balanced accuracy increases from 36.1% for H3 to 45.6% after routing, reflecting class-specific recall

Table 2: Task-Level Qualification Outcomes

Candidate	Qualification evidence	n	Observed result	Authority decision
Futures positioning specialist	Frozen weekly confirmation	104	Rank quality .568 vs .433, 2013/2014 values .616/.520, 5/5 offsets > 0.5, interval score 42.53 vs 61.49	Forecast authority granted
Gasoline CPI specialist	Unused monthly holdout	20	MAE 2.252 vs 2.701, 16.6% lower MAE	Gasoline forecast authority granted
Earnings-per-share (EPS) growth upgrade	Frozen confirmation	32	Mean quality-score change -0.016	Base path retained
Macro-revision candidate	Historical holdout	18	50.0% accuracy	LLM path retained
Treasury auction candidate	Interval replay	7	89.3% leave-one-tenor-out coverage	Existing point path retained
Yield-curve candidate	Surrogate stress test	19	No real historical multi-draw archive	LLM path retained
Issuer-credit candidate	Historical retrieval replay	2	Two historical retrieval cases	LLM path retained
EPS year-over-year selector	Retrospective real-LLM replay	27	H3 and selected accuracy 70.4%, three-draw unanimity 100%	Classification accuracy unchanged

Notes: The n column reports each task’s evaluation units, including weekly origins for futures and tenors for Treasury auctions. Treasury coverage pools 28 held-out predictions across seven tenor groups. Rows outside futures positioning and gasoline are secondary scope audits of task-specific gates or evidence sufficiency. Rank quality equals $(1 + \rho_S)/2$, where ρ_S is the within-origin Spearman rank correlation. Lower proper interval scores and lower gasoline MAE indicate better forecast performance. Each task uses its own frozen control during specialist qualification.

in the imbalanced sample. Nominal 90% intervals cover 23.3% of outcomes under H3 and 26.7% after routing, below the nominal level in both cases.

The preregistered Gemini 3.5 Flash check yields no paired accuracy change, and its intervention set shares no events with the primary run. The exploratory Gemini 3.1 Pro check yields a change of -6.7 percentage points, with a 95% interval of $[-23.3, 10.0]$ (Table 3). The same-event checks produce different intervention sets and accuracy changes, with all three confidence intervals spanning zero.

3.3 Record Preservation After Authority Transfer

Figure 2A shows where routing activates in the primary outcome-sealed Gemini 3.8 run. Twenty-three of the 30 events have unanimous three-draw labels, and three more have a two-to-one split with the auxiliary classifier on the majority class. All 26 of those cases retain the complete H3 baseline record. The remaining four events satisfy the authority-transfer rule, representing 13.3% of the primary sample. The auxiliary prediction takes the previously unchosen third class in three


Table 3: Backbone Dependence of Post-Earnings Routing

Backbone	Evidence status	H3 accuracy	Routed accuracy	Change [95% confidence interval]	McNe-mar p	H3 / routed coverage	Interventions
Gemini 3.8 Flash	Outcome-sealed historical holdout	26.7%	33.3%	+6.7 [−6.7, +20.0]	0.625	23.3% / 26.7%	4: 3 / 1 / 0
Gemini 3.5 Flash	Same-event preregistered check	26.7%	26.7%	0.0 [−10.0, +10.0]	1.000	23.3% / 23.3%	7: 1 / 1 / 5
Gemini 3.1 Pro	Same-event exploratory check	40.0%	33.3%	−6.7 [−23.3, +10.0]	0.688	26.7% / 30.0%	9: 2 / 4 / 3

Notes: H3 is the three-draw majority baseline used for each backbone comparison. The Gemini 3.8 Flash results come from the primary outcome-sealed historical evaluation. Gemini 3.5 Flash and Gemini 3.1 Pro reuse the same events after outcome release. The reported confidence intervals use bootstrap resampling at the issuer-event level. All changes in forecast accuracy are reported in percentage points. The coverage column reports H3 first and the routed forecast second. Intervention counts report total interventions followed by corrected, harmed, and unchanged cases.

events and the one-vote minority class in one. The two rightmost states in Figure 2A contain every intervention.

Figure 2: Post-Earnings Authority Transfer



Notes: Panel A reports H3 and routed accuracy by disagreement state, with event counts and routing rates below each state. The two rightmost states in Panel A contain all four interventions. Panel B compares point-class consistency under complete-record finalization and the class-only patch.

Figure 2B reports the structural-integrity check on the same four category changes. Complete-record finalization preserves point-class consistency in all four routed cases. The class-only patch uses the same selected classes, so classification accuracy is identical under the two update rules. The class-only patch leaves all four changed records inconsistent with their numerical point estimates.

All 30 records produced by the routing rule remain structurally valid with exact evidence-quote matching.

4 Discussion

The qualification results tie numerical specialist authority to a validated financial target, input contract, and historical evidence base. Backbone changes alter which cases trigger routing even when the decision rule remains fixed, so evaluation also needs to cover the routing configuration (Sculley et al., 2015; Ribeiro et al., 2020). The observed changes in intervention sets motivate evaluation of the target, backbone, specialist, and routing rule as one forecasting configuration.

Complete-record finalization links the selected category to the numerical fields and evidence quote supplied to analysts or risk systems. Record coherence and calibrated uncertainty remain separate properties, and the low post-earnings coverage constrains probability-sensitive uses. Tail-risk decisions in prediction markets also depend on forecast calibration (Gong & Dimitrov, 2026). The primary evidence covers 30 events, and the later backbone checks reuse those events after outcome release, limiting conclusions about new market conditions. Anonymization and outcome sealing control information access during evaluation, but do not establish that the evaluated events were absent from model training.

5 Conclusion

Task-specific qualification and complete-record finalization govern different stages of financial prediction, from selecting a forecaster to updating the record after a category changes. The evaluated specialists meet their numerical gates, and the routing checks identify configuration-sensitive interventions with consistent reconstructed records. Rolling evaluations across new events and model versions should test whether requalification preserves specialist performance and whether routing improves forecasts beyond the existing controls.

References

- Bates, J. M., & Granger, C. W. J. (1969) The Combination of Forecasts. *Operational Research Quarterly*, 20(4), 451–468.
- Brown, S. J., & Warner, J. B. (1985) Using Daily Stock Returns: The Case of Event Studies.

- Journal of Financial Economics*, 14(1), 3–31.
- Commodity Futures Trading Commission. (2026) Historical Compressed Commitments of Traders Reports. <https://www.cftc.gov/MarketReports/CommitmentsofTraders/HistoricalCompressed/index.htm>. Accessed September 29, 2026.
- Chow, C. K. (1970) On Optimum Recognition Error and Reject Tradeoff. *IEEE Transactions on Information Theory*, 16(1), 41–46.
- Clemen, R. T. (1989) Combining Forecasts: A Review and Annotated Bibliography. *International Journal of Forecasting*, 5(4), 559–583.
- Croushore, D., & Stark, T. (2001) A Real-Time Data Set for Macroeconomists. *Journal of Econometrics*, 105(1), 111–130.
- Dawid, A. P. (1984) Present Position and Potential Developments: Some Personal Views: Statistical Theory: The Prequential Approach. *Journal of the Royal Statistical Society, Series A*, 147(2), 278–292.
- Farquhar, S., Kossen, J., Kuhn, L., & Gal, Y. (2024) Detecting hallucinations in large language models using semantic entropy. *Nature*, 630, 625–630.
- Geifman, Y., & El-Yaniv, R. (2017) Selective Classification for Deep Neural Networks. *Advances in Neural Information Processing Systems*, 30.
- Gneiting, T., Balabdaoui, F., & Raftery, A. E. (2007) Probabilistic Forecasts, Calibration and Sharpness. *Journal of the Royal Statistical Society, Series B*, 69(2), 243–268.
- Gneiting, T., & Raftery, A. E. (2007) Strictly Proper Scoring Rules, Prediction, and Estimation. *Journal of the American Statistical Association*, 102(477), 359–378.
- Gong, G., & Dimitrov, S. (2026) When Text Helps: AI Text Signals for Calibration-Sensitive Tail-Risk Decisions in Prediction Markets. *SSRN Working Paper*.
- Google DeepMind. (2026) Gemini 3.8 Flash: Model Card. <https://deepmind.google/models/model-cards/gemini-3-8-flash/>. Published September 2, 2026. Accessed October 7, 2026.
- Hyndman, R. J., & Koehler, A. B. (2006) Another Look at Measures of Forecast Accuracy. *International Journal of Forecasting*, 22(4), 679–688.
- Jin, W., Khanna, R., Kim, S., Lee, D.-H., Morstatter, F., Galstyan, A., & Ren, X. (2021) ForecastQA: A Question Answering Challenge for Event Forecasting with Temporal Text Data. *Proceedings of ACL-IJCNLP 2021*, 4636–4650.
- Kuleshov, V., Fenner, N., & Ermon, S. (2018) Accurate Uncertainties for Deep Learning Using Calibrated Regression. *Proceedings of the 35th International Conference on Machine Learning*,

80, 2796–2804.

- Lewis, P., et al. (2020) Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. *Advances in Neural Information Processing Systems*, 33.
- MacKinlay, A. C. (1997) Event Studies in Economics and Finance. *Journal of Economic Literature*, 35(1), 13–39.
- Madras, D., Pitassi, T., & Zemel, R. (2018) Predict Responsibly: Improving Fairness and Accuracy by Learning to Defer. *Advances in Neural Information Processing Systems*, 31.
- Manakul, P., Liusie, A., & Gales, M. (2023) SelfCheckGPT: Zero-Resource Black-Box Hallucination Detection for Generative Large Language Models. *Proceedings of EMNLP 2023*, 9004–9017.
- Mozannar, H., & Sontag, D. (2020) Consistent Estimators for Learning to Defer to an Expert. *Proceedings of the 37th International Conference on Machine Learning*, 119, 7076–7087.
- New York Stock Exchange. (2026) 2026 Annual Guidance Letter. https://www.nyse.com/publicdocs/nyse/markets/nyse/NYSE_2026_Annual_Guidance_Letter.pdf. Accessed September 29, 2026.
- Ribeiro, M. T., Wu, T., Guestrin, C., & Singh, S. (2020) Beyond Accuracy: Behavioral Testing of NLP Models with CheckList. *Proceedings of ACL 2020*, 4902–4912.
- Sculley, D., et al. (2015) Hidden Technical Debt in Machine Learning Systems. *Advances in Neural Information Processing Systems*, 28.
- U.S. Bureau of Labor Statistics. (2025) Handbook of Methods: Consumer Price Index—Overview. <https://www.bls.gov/opub/hom/cpi/>. Accessed September 29, 2026.
- U.S. Energy Information Administration. (2026) Methodology for EIA Weekly Retail Gasoline Price Estimates. https://www.eia.gov/petroleum/gasdiesel/gas_proc-methods.php. Accessed September 29, 2026.
- U.S. Securities and Exchange Commission. (2025) Form 8-K: Current Report Pursuant to Section 13 or 15(d). <https://www.sec.gov/submit-filings/forms-index>. Accessed September 29, 2026.
- Yahoo Finance. (2026) What Is the Adjusted Close? <https://help.yahoo.com/kb/SLN28256.html>. Accessed September 29, 2026.